

STAT 528 - Advanced Regression Analysis II

Multivariate Regression

Daniel J. Eck
Department of Statistics
University of Illinois

Learning Objectives Today

- ▶ Multivariate regression modeling
- ▶ Examples

Motor Trend Cars example

The data was extracted from the 1974 Motor Trend US magazine, and comprises fuel consumption and 10 aspects of automobile design and performance for 32 automobiles (1973-74 models). The variables are:

- ▶ mpg: Miles/(US) gallon
- ▶ disp: Displacement (cu.in.)
- ▶ hp: Gross horsepower
- ▶ wt: Weight (1000 lbs)
- ▶ cyl: Number of cylinders
- ▶ am: Transmission (0 = automatic, 1 = manual)
- ▶ carb: Number of carburetors

The first four variables are response variables corresponding to engine performance and size. The next three variables are engine design variables.

We load in the data

```
data(mtcars)
head(mtcars)
```

```
##           mpg  cyl  disp  hp  drat    wt  qsec vs  am  gear  carb
## Mazda RX4      21.0   6  160  110  3.90  2.620  16.46  0  1    4    4
## Mazda RX4 Wag  21.0   6  160  110  3.90  2.875  17.02  0  1    4    4
## Datsun 710     22.8   4  108   93  3.85  2.320  18.61  1  1    4    1
## Hornet 4 Drive  21.4   6  258  110  3.08  3.215  19.44  1  0    3    1
## Hornet Sportabout 18.7   8  360  175  3.15  3.440  17.02  0  0    3    2
## Valiant        18.1   6  225  105  2.76  3.460  20.22  1  0    3    1
```

The standard `lm` function in R can fit multivariate linear regression models.

```
mtcars$cyl <- factor(mtcars$cyl)
Y <- as.matrix(mtcars[,c("mpg", "disp", "hp", "wt")])
m <- lm(Y ~ cyl + am + carb, data=mtcars, x = TRUE)
```

```
# estimate of beta'
betahat <- coef(m)
betahat
```

```
##           mpg      disp      hp      wt
## (Intercept) 25.320303 134.32487 46.5201421 2.7612069
## cyl6       -3.549419  61.84324  0.9116288  0.1957229
## cyl8       -6.904637 218.99063 87.5910956  0.7723077
## am         4.226774 -43.80256  4.4472569 -1.0254749
## carb      -1.119855  1.72629 21.2764930  0.1749132
```

We now estimate Σ via MLE and provide an unbiased estimator.

```
# estimates of Sigma
SSE <- crossprod(Y - m$fitted.values)
n <- nrow(Y)
p <- nrow(coef(m))
SigmaMLE <- SSE / n
SigmaMLE
```

```
##           mpg      disp          hp      wt
## mpg      6.638633 -44.94796 -16.6232233 -0.5548030
## disp -44.947964 2113.48487 358.7058833 15.2773945
## hp    -16.623223 358.70588 487.0718440 0.3933977
## wt    -0.554803 15.27739 0.3933977 0.2171394
```

```
Sigmahat <- SSE / (n - p)
Sigmahat
```

```
##           mpg      disp          hp      wt
## mpg      7.8680094 -53.27166 -19.7015979 -0.6575443
## disp -53.2716607 2504.87095 425.1328988 18.1065416
## hp    -19.7015979 425.13290 577.2703337 0.4662491
## wt    -0.6575443 18.10654 0.4662491 0.2573503
```

We can see that the R's `vcov` function provides an estimate of $\text{Var}(\text{vec}(\hat{\beta}')) = \Sigma \otimes (\mathbf{X}'\mathbf{X})^{-1}$ as its default

```
X <- m$x  
unique(round(vcov(m) - kronecker(Sigmahat, solve(crossprod(X))), 10))
```

```
##                mpg:(Intercept) mpg:cyl6 mpg:cyl8 mpg:am mpg:carb  
## mpg:(Intercept)                0         0         0         0  
##                disp:(Intercept) disp:cyl6 disp:cyl8 disp:am disp:carb  
## mpg:(Intercept)                0         0         0         0  
##                hp:(Intercept)  hp:cyl6  hp:cyl8 hp:am  hp:carb wt:(Intercept)  
## mpg:(Intercept)                0         0         0         0  
##                wt:cyl6 wt:cyl8 wt:am wt:carb  
## mpg:(Intercept)                0         0         0
```

We obtain inferences for regression coefficients corresponding to the first response variable mpg.

```
# summary table from lm
```

```
msum <- summary(m)
```

```
msum[[1]]
```

```
##
```

```
## Call:
```

```
## lm(formula = mpg ~ cyl + am + carb, data = mtcars, x = TRUE)
```

```
##
```

```
## Residuals:
```

```
##      Min       1Q   Median       3Q      Max  
## -5.9074 -1.1723  0.2538  1.4851  5.4728
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error t value Pr(>|t|)  
## (Intercept)  25.3203      1.2238  20.690 < 2e-16 ***  
## cyl6         -3.5494      1.7296  -2.052 0.049959 *  
## cyl8         -6.9046      1.8078  -3.819 0.000712 ***  
## am           4.2268      1.3499   3.131 0.004156 **  
## carb        -1.1199      0.4354  -2.572 0.015923 *
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
```

```
## Residual standard error: 2.805 on 27 degrees of freedom
```

```
## Multiple R-squared:  0.8113, Adjusted R-squared:  0.7834
```

```
## F-statistic: 29.03 on 4 and 27 DF,  p-value: 1.991e-09
```


We compare the summary table on the previous design to one obtained directly from theory

```
## estimates and standard errors
msum2 <- cbind(coef(m)[, 1], sqrt(diag( kronecker(Sigmahat, solve(crossprod(X))) ))[1:5])

## tstat and p-values
msum2 <- cbind(msum2, msum2[, 1] / msum2[, 2])
msum2 <- cbind(msum2, sapply(msum2[, 3], function(x) pt(abs(x), df = n - p, lower = FALSE)*2 ))
msum2
```

```
##           [,1]      [,2]      [,3]      [,4]
## (Intercept) 25.320303 1.2237903 20.690067 4.303483e-18
## cyl6        -3.549419 1.7295506 -2.052221 4.995947e-02
## cyl8        -6.904637 1.8078219 -3.819313 7.124146e-04
## am           4.226774 1.3499249  3.131118 4.156214e-03
## carb        -1.119855 0.4353558 -2.572274 1.592280e-02
```

Inference and nested models

There are a plethora of additional tests that we could perform in this setting (we will not stress each test's origins in this course).

These include the Wilk's lambda, Pillai's trace, Hotelling-Lawley trace, and Roy's greatest root.

First, let $E = n\tilde{\Sigma}$ where $\tilde{\Sigma}$ is the MLE for the full model with β unconstrained and let $\tilde{H} = n(\tilde{\Sigma}_1 - \tilde{\Sigma})$ where $\tilde{\Sigma}_1$ is the MLE of Σ in the reduced model constrained by $\beta_2 = 0$.

The test statistics now follow:

- ▶ Wilk's lambda $= \prod_{i=1}^s \frac{1}{1+\eta_i} = \frac{|\tilde{E}|}{|\tilde{E}+\tilde{H}|}$
- ▶ Pillai's trace $= \sum_{i=1}^s \frac{\eta_i}{1+\eta_i} = \text{tr}[\tilde{H}(\tilde{E} + \tilde{H})^{-1}]$
- ▶ Hotelling-Lawley trace $= \sum_{i=1}^s \eta_i = \text{tr}(\tilde{H}\tilde{E}^{-1})$
- ▶ Roy's greatest root $= \frac{\eta_1}{1+\eta_1}$

where $\eta_1 \geq \eta_2 \geq \dots \geq \eta_s$ denote the nonzero eigenvalues of $\tilde{H}\tilde{E}^{-1}$.

We demonstrate estimation of Wilk's lambda and Pillai's trace.

```
## anova implements these methods
m0 <- lm(Y ~ am + carb, data=mtcars)
anova(m0, m, test="Wilks")

## Analysis of Variance Table
##
## Model 1: Y ~ am + carb
## Model 2: Y ~ cyl + am + carb
##   Res.Df Df Gen.var.   Wilks approx F num Df den Df    Pr(>F)
## 1      29      43.692
## 2      27 -2    29.862 0.16395    8.8181      8    48 2.525e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
anova(m0, m, test="Pillai")
```

```
## Analysis of Variance Table
##
## Model 1: Y ~ am + carb
## Model 2: Y ~ cyl + am + carb
##   Res.Df Df Gen.var. Pillai approx F num Df den Df    Pr(>F)
## 1      29      43.692
## 2      27 -2    29.862 1.0323    6.6672      8    50 6.593e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The same quantities are estimated directly

```
Etilde <- n * SigmaMLE
SigmaTilde1 <- crossprod(Y - m0$fitted.values) / n
Htilde <- n * (SigmaTilde1 - SigmaMLE)
HEi <- Htilde %*% solve(Etilde)
HEi.values <- eigen(HEi)$values
c(Wilks = prod(1 / (1 + HEi.values)), Pillai = sum(HEi.values / (1 + HEi.values)))
```

```
##           Wilks           Pillai
## 0.1639527+0i 1.0322975+0i
```

We will now demonstrate confidence and prediction intervals in the motor trends cars example. Note that we have to code our own functions because R does not have the capability to provide these quantities. R does provide point predictions.

```
# confidence interval
newdata <- data.frame(cyl=factor(6, levels=c(4,6,8)), am=1, carb=4)
predict(m, newdata, interval="confidence")
```

```
##      mpg      disp      hp      wt
## 1 21.51824 159.2707 136.985 2.631108
```

```
# prediction interval
newdata <- data.frame(cyl=factor(6, levels=c(4,6,8)), am=1, carb=4)
predict(m, newdata, interval="prediction")
```

```
##      mpg      disp      hp      wt
## 1 21.51824 159.2707 136.985 2.631108
```

Here is the function which produces confidence and prediction intervals (credit to Nathaniel Helwig)

We now provide 95% confidence and prediction intervals using code from [Nathaniel Helwig](#) (see Rmd file):

```
# confidence interval
newdata <- data.frame(cyl=factor(6, levels=c(4,6,8)), am=1, carb=4)
pred.mlm(m, newdata)
```

```
##           mpg      disp          hp      wt
## fit 21.51824 159.2707 136.98500 2.631108
## lwr 16.65593  72.5141  95.33649 1.751736
## upr 26.38055 246.0273 178.63351 3.510479
```

```
# prediction interval
newdata <- data.frame(cyl=factor(6, levels=c(4,6,8)), am=1, carb=4)
pred.mlm(m, newdata, interval="prediction")
```

```
##           mpg      disp          hp      wt
## fit 21.518240 159.27070 136.98500 2.6311076
## lwr  9.680053 -51.95435  35.58397 0.4901152
## upr 33.356426 370.49576 238.38603 4.7720999
```